

Can a model detect what we change inside it?

We perturb internal activations, then ask the model what happened.

EXPERIMENT 1

Where was the perturbation?

A The cat sat on the mat.

← MODIFIED

B The door was painted blue.

Intervention on A's activations in this example

"Which sentence was targeted?"


Forced choice: A or B

HYPOTHESES H1

The choice follows the injection site. Detection differs across perturbation types. Normalizing dose changes these differences.

MULTI-INJECTION / EXPERIMENTS 7-9

How much can it report?



"How many?" 0, 1, 2, 3 or 4

H7 Counting reflects more than total intensity.

"Which concepts?" Select two labels

H8 The model identifies both injected concepts.

"Which was earlier?" Choose A or B

H9 The model identifies the earlier injection.